

# Deep Residual Learning for Image Recognition

Project Members:

- Aditya Pradeep Waghmode - Z23737910
- Surendra Ramisetty - Z23734231
- Srilatha Jidugu - Z23747421

## Trail 1: Code

**Google Colab link for trail 1 code:**

<https://colab.research.google.com/drive/1fKeMZVo7Rw2LL6zvBPTy9GEzH1mzPXGn?usp=sharing>

▼ Import necessary libraries

```
✓ [4] import numpy as np
4s from keras.datasets import cifar10
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, Activation
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
```

▼ Load CIFAR-10 dataset and split into training and testing sets

```
✓ [5] (x_train, y_train), (x_test, y_test) = cifar10.load_data()
6s
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 4s 0us/step
```

---

- ▼ Create a validation set by randomly selecting 20% of the training images

```
✓ [6] validation_split = 0.2  
0s validation_samples = int(len(x_train) * validation_split)  
x_val = x_train[:validation_samples]  
y_val = y_train[:validation_samples]  
x_train = x_train[validation_samples:]  
y_train = y_train[validation_samples:]
```

- ▼ Scale pixel values to a range between 0 and 1

```
✓ [7] x_train = x_train.astype('float32') / 255  
0s x_val = x_val.astype('float32') / 255  
x_test = x_test.astype('float32') / 255
```

- ▼ Convert labels to binary class matrices

```
✓ [8] y_train = to_categorical(y_train, 10)  
s y_val = to_categorical(y_val, 10)  
y_test = to_categorical(y_test, 10)
```

Double-click (or enter) to edit

- ▼ Import necessary libraries for model creation

```
✓ [9] import tensorflow as tf  
s from tensorflow import keras  
from tensorflow.keras import layers  
from tensorflow.keras.layers import Dense, Flatten, Activation  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.optimizers import Adam  
from keras.preprocessing.image import ImageDataGenerator  
from keras.src.layers.serialization import activation
```

- ▼ Build CNN model with data augmentation

```
✓ [10] data_generator = ImageDataGenerator(rotation_range=10, width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)  
50m data_generator.fit(x_train)  
  
model_augmented = Sequential()  
pretraining_model = tf.keras.applications.ResNet152V2(  
    include_top=False,  
    weights="imagenet",  
    input_tensor=None,  
    input_shape=(32, 32, 3),  
    pooling=None, # No global pooling in the pre-trained model  
    classes=10  
)  
  
for layer in pretraining_model.layers:  
    layer.trainable = True
```

```

model_augmented.add(pretraining_model)
model_augmented.add(Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(32, 32, 3)))
model_augmented.add(Activation('relu'))
model_augmented.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model_augmented.add(Activation('relu'))
model_augmented.add(MaxPooling2D(pool_size=(1, 1)))
model_augmented.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model_augmented.add(Activation('relu'))
model_augmented.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model_augmented.add(Activation('relu'))
model_augmented.add(MaxPooling2D(pool_size=(1, 1)))
model_augmented.add(Flatten())
model_augmented.add(Dense(512, activation='relu'))
model_augmented.add(Dropout(0.5))
model_augmented.add(Dense(10, activation='softmax'))

model_augmented.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.001), metrics=['accuracy'])
checkpoint_augmented = ModelCheckpoint('best_model_augmented.h5', save_best_only=True, monitor='val_loss', mode='min', verbose=1)
history_augmented = model_augmented.fit(data_generator.flow(x_train, y_train, batch_size=512), epochs=50, validation_data=(x_val, y_val), callbacks=[checkpoint_augmented])
model_augmented.load_weights('best_model_augmented.h5')
train_loss_augmented, train_accuracy_augmented = model_augmented.evaluate(x_train, y_train)
val_loss_augmented, val_accuracy_augmented = model_augmented.evaluate(x_val, y_val)

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5)

234545216/234545216 [=====] - 1s 0us/step

WARNING:absl: `lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g., `tf.keras.optimizers.legacy.Adam`.

Epoch 1/50

79/79 [=====] - ETA: 0s - loss: 2.2148 - accuracy: 0.1395

Epoch 1: val\_loss improved from inf to 1283.32739, saving model to best\_model\_augmented.h5

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3079: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')`.

saving\_api.save\_model(

79/79 [=====] - 159s 654ms/step - loss: 2.2148 - accuracy: 0.1395 - val\_loss: 1283.3274 - val\_accuracy: 0.0974

Epoch 2/50

79/79 [=====] - ETA: 0s - loss: 1.8325 - accuracy: 0.2298

Epoch 2: val\_loss improved from 1283.32739 to 3.86596, saving model to best\_model\_augmented.h5

79/79 [=====] - 59s 748ms/step - loss: 1.8325 - accuracy: 0.2298 - val\_loss: 3.8660 - val\_accuracy: 0.1880

Epoch 3/50

79/79 [=====] - ETA: 0s - loss: 1.6414 - accuracy: 0.3225

Epoch 3: val\_loss did not improve from 3.86596

79/79 [=====] - 45s 564ms/step - loss: 1.6414 - accuracy: 0.3225 - val\_loss: 7.4576 - val\_accuracy: 0.2724

Epoch 4/50

79/79 [=====] - ETA: 0s - loss: 1.4701 - accuracy: 0.4108

Epoch 4: val\_loss did not improve from 3.86596

79/79 [=====] - 42s 528ms/step - loss: 1.4701 - accuracy: 0.4108 - val\_loss: 3.9650 - val\_accuracy: 0.2928

Epoch 5/50

79/79 [=====] - ETA: 0s - loss: 1.4853 - accuracy: 0.4238

Epoch 5: val\_loss improved from 3.86596 to 2.14562, saving model to best\_model\_augmented.h5

79/79 [=====] - 51s 642ms/step - loss: 1.4853 - accuracy: 0.4238 - val\_loss: 2.1456 - val\_accuracy: 0.4281

Epoch 6/50

79/79 [=====] - ETA: 0s - loss: 1.4107 - accuracy: 0.4665

Epoch 6: val\_loss did not improve from 2.14562

79/79 [=====] - 45s 564ms/step - loss: 1.4107 - accuracy: 0.4665 - val\_loss: 2.2325 - val\_accuracy: 0.4260

Epoch 7/50

79/79 [=====] - ETA: 0s - loss: 1.3240 - accuracy: 0.5099

Epoch 7: val\_loss did not improve from 2.14562

79/79 [=====] - 44s 551ms/step - loss: 1.3240 - accuracy: 0.5099 - val\_loss: 2.4001 - val\_accuracy: 0.3953  
Epoch 8/50  
79/79 [=====] - ETA: 0s - loss: 1.1981 - accuracy: 0.5767  
Epoch 8: val\_loss improved from 2.14562 to 1.70583, saving model to best\_model\_augmented.h5  
79/79 [=====] - 52s 655ms/step - loss: 1.1981 - accuracy: 0.5767 - val\_loss: 1.7058 - val\_accuracy: 0.5855  
Epoch 9/50  
79/79 [=====] - ETA: 0s - loss: 1.0740 - accuracy: 0.6270  
Epoch 9: val\_loss improved from 1.70583 to 1.16933, saving model to best\_model\_augmented.h5  
79/79 [=====] - 52s 652ms/step - loss: 1.0740 - accuracy: 0.6270 - val\_loss: 1.1693 - val\_accuracy: 0.6300  
Epoch 10/50  
79/79 [=====] - ETA: 0s - loss: 0.9996 - accuracy: 0.6625  
Epoch 10: val\_loss improved from 1.16933 to 1.15764, saving model to best\_model\_augmented.h5  
79/79 [=====] - 52s 652ms/step - loss: 0.9996 - accuracy: 0.6625 - val\_loss: 1.1576 - val\_accuracy: 0.6614  
Epoch 11/50  
79/79 [=====] - ETA: 0s - loss: 0.9315 - accuracy: 0.6896  
Epoch 11: val\_loss improved from 1.15764 to 1.04676, saving model to best\_model\_augmented.h5  
79/79 [=====] - 49s 625ms/step - loss: 0.9315 - accuracy: 0.6896 - val\_loss: 1.0468 - val\_accuracy: 0.6617  
Epoch 12/50  
79/79 [=====] - ETA: 0s - loss: 0.8714 - accuracy: 0.7095  
Epoch 12: val\_loss improved from 1.04676 to 0.97995, saving model to best\_model\_augmented.h5  
79/79 [=====] - 53s 674ms/step - loss: 0.8714 - accuracy: 0.7095 - val\_loss: 0.9799 - val\_accuracy: 0.6934  
Epoch 13/50  
79/79 [=====] - ETA: 0s - loss: 0.8218 - accuracy: 0.7238  
Epoch 13: val\_loss did not improve from 0.97995  
79/79 [=====] - 42s 529ms/step - loss: 0.8218 - accuracy: 0.7238 - val\_loss: 1.2358 - val\_accuracy: 0.6334  
Epoch 14/50  
79/79 [=====] - ETA: 0s - loss: 0.7795 - accuracy: 0.7421  
Epoch 14: val\_loss improved from 0.97995 to 0.78946, saving model to best\_model\_augmented.h5  
79/79 [=====] - 55s 699ms/step - loss: 0.7795 - accuracy: 0.7421 - val\_loss: 0.7895 - val\_accuracy: 0.7464  
Epoch 15/50  
79/79 [=====] - ETA: 0s - loss: 0.7563 - accuracy: 0.7521  
Epoch 15: val\_loss did not improve from 0.78946  
79/79 [=====] - 44s 556ms/step - loss: 0.7563 - accuracy: 0.7521 - val\_loss: 1.0822 - val\_accuracy: 0.6820  
Epoch 16/50  
79/79 [=====] - ETA: 0s - loss: 0.7327 - accuracy: 0.7608  
Epoch 16: val\_loss improved from 0.78946 to 0.76350, saving model to best\_model\_augmented.h5  
79/79 [=====] - 50s 631ms/step - loss: 0.7327 - accuracy: 0.7608 - val\_loss: 0.7635 - val\_accuracy: 0.7566  
Epoch 17/50  
79/79 [=====] - ETA: 0s - loss: 0.7046 - accuracy: 0.7702  
Epoch 17: val\_loss did not improve from 0.76350  
79/79 [=====] - 44s 552ms/step - loss: 0.7046 - accuracy: 0.7702 - val\_loss: 0.8856 - val\_accuracy: 0.7191  
Epoch 18/50  
79/79 [=====] - ETA: 0s - loss: 0.6594 - accuracy: 0.7856  
Epoch 18: val\_loss did not improve from 0.76350

79/79 [=====] - 45s 566ms/step - loss: 0.6594 - accuracy: 0.7856 - val\_loss: 0.8845 - val\_accuracy: 0.7286  
Epoch 19/50  
79/79 [=====] - ETA: 0s - loss: 0.6426 - accuracy: 0.7910  
Epoch 19: val\_loss did not improve from 0.76350  
79/79 [=====] - 43s 542ms/step - loss: 0.6426 - accuracy: 0.7910 - val\_loss: 0.9594 - val\_accuracy: 0.7147  
Epoch 20/50  
79/79 [=====] - ETA: 0s - loss: 0.6377 - accuracy: 0.7928  
Epoch 20: val\_loss did not improve from 0.76350  
79/79 [=====] - 44s 556ms/step - loss: 0.6377 - accuracy: 0.7928 - val\_loss: 0.8898 - val\_accuracy: 0.7321  
Epoch 21/50  
79/79 [=====] - ETA: 0s - loss: 0.6053 - accuracy: 0.8020  
Epoch 21: val\_loss did not improve from 0.76350  
79/79 [=====] - 42s 532ms/step - loss: 0.6053 - accuracy: 0.8020 - val\_loss: 0.8557 - val\_accuracy: 0.7541  
Epoch 22/50  
79/79 [=====] - ETA: 0s - loss: 0.5857 - accuracy: 0.8103  
Epoch 22: val\_loss improved from 0.76350 to 0.73701, saving model to best\_model\_augmented.h5  
79/79 [=====] - 55s 700ms/step - loss: 0.5857 - accuracy: 0.8103 - val\_loss: 0.7370 - val\_accuracy: 0.7733  
Epoch 23/50  
79/79 [=====] - ETA: 0s - loss: 0.5693 - accuracy: 0.8164  
Epoch 23: val\_loss did not improve from 0.73701  
79/79 [=====] - 41s 521ms/step - loss: 0.5693 - accuracy: 0.8164 - val\_loss: 0.7746 - val\_accuracy: 0.7581  
Epoch 24/50  
79/79 [=====] - ETA: 0s - loss: 0.5504 - accuracy: 0.8238  
Epoch 24: val\_loss did not improve from 0.73701  
79/79 [=====] - 44s 558ms/step - loss: 0.5504 - accuracy: 0.8238 - val\_loss: 0.8447 - val\_accuracy: 0.7334  
Epoch 25/50  
79/79 [=====] - ETA: 0s - loss: 0.5465 - accuracy: 0.8206  
Epoch 25: val\_loss did not improve from 0.73701  
79/79 [=====] - 45s 566ms/step - loss: 0.5465 - accuracy: 0.8206 - val\_loss: 0.8152 - val\_accuracy: 0.7591  
Epoch 26/50  
79/79 [=====] - ETA: 0s - loss: 0.5265 - accuracy: 0.8277  
Epoch 26: val\_loss improved from 0.73701 to 0.66414, saving model to best\_model\_augmented.h5  
79/79 [=====] - 53s 670ms/step - loss: 0.5265 - accuracy: 0.8277 - val\_loss: 0.6641 - val\_accuracy: 0.7972  
Epoch 27/50  
79/79 [=====] - ETA: 0s - loss: 0.5086 - accuracy: 0.8386  
Epoch 27: val\_loss did not improve from 0.66414  
79/79 [=====] - 44s 552ms/step - loss: 0.5086 - accuracy: 0.8386 - val\_loss: 0.7073 - val\_accuracy: 0.7749  
Epoch 28/50  
79/79 [=====] - ETA: 0s - loss: 0.4855 - accuracy: 0.8437  
Epoch 28: val\_loss did not improve from 0.66414  
79/79 [=====] - 42s 523ms/step - loss: 0.4855 - accuracy: 0.8437 - val\_loss: 1.5998 - val\_accuracy: 0.6233  
Epoch 29/50  
79/79 [=====] - ETA: 0s - loss: 0.4977 - accuracy: 0.8379  
Epoch 29: val\_loss improved from 0.66414 to 0.65613, saving model to best\_model\_augmented.h5

79/79 [=====] - 51s 650ms/step - loss: 0.4977 - accuracy: 0.8379 - val\_loss: 0.6561 - val\_accuracy: 0.7943  
Epoch 30/50  
79/79 [=====] - ETA: 0s - loss: 0.4630 - accuracy: 0.8511  
Epoch 30: val\_loss did not improve from 0.65613  
79/79 [=====] - 43s 548ms/step - loss: 0.4630 - accuracy: 0.8511 - val\_loss: 0.8926 - val\_accuracy: 0.7512  
Epoch 31/50  
79/79 [=====] - ETA: 0s - loss: 0.4605 - accuracy: 0.8510  
Epoch 31: val\_loss did not improve from 0.65613  
79/79 [=====] - 44s 550ms/step - loss: 0.4605 - accuracy: 0.8510 - val\_loss: 0.9711 - val\_accuracy: 0.7432  
Epoch 32/50  
79/79 [=====] - ETA: 0s - loss: 0.4621 - accuracy: 0.8508  
Epoch 32: val\_loss did not improve from 0.65613  
79/79 [=====] - 44s 551ms/step - loss: 0.4621 - accuracy: 0.8508 - val\_loss: 0.7870 - val\_accuracy: 0.7800  
Epoch 33/50  
79/79 [=====] - ETA: 0s - loss: 0.4504 - accuracy: 0.8537  
Epoch 33: val\_loss did not improve from 0.65613  
79/79 [=====] - 44s 557ms/step - loss: 0.4504 - accuracy: 0.8537 - val\_loss: 0.8519 - val\_accuracy: 0.7504  
Epoch 34/50  
79/79 [=====] - ETA: 0s - loss: 0.4414 - accuracy: 0.8558  
Epoch 34: val\_loss did not improve from 0.65613  
79/79 [=====] - 44s 553ms/step - loss: 0.4414 - accuracy: 0.8558 - val\_loss: 0.6652 - val\_accuracy: 0.7952  
Epoch 35/50  
79/79 [=====] - ETA: 0s - loss: 0.4252 - accuracy: 0.8618  
Epoch 35: val\_loss did not improve from 0.65613  
79/79 [=====] - 43s 536ms/step - loss: 0.4252 - accuracy: 0.8618 - val\_loss: 0.7694 - val\_accuracy: 0.7682  
Epoch 36/50  
79/79 [=====] - ETA: 0s - loss: 0.4058 - accuracy: 0.8674  
Epoch 36: val\_loss did not improve from 0.65613  
79/79 [=====] - 43s 536ms/step - loss: 0.4058 - accuracy: 0.8674 - val\_loss: 0.6575 - val\_accuracy: 0.8081  
Epoch 37/50  
79/79 [=====] - ETA: 0s - loss: 0.3974 - accuracy: 0.8726  
Epoch 37: val\_loss did not improve from 0.65613  
79/79 [=====] - 43s 540ms/step - loss: 0.3974 - accuracy: 0.8726 - val\_loss: 0.6886 - val\_accuracy: 0.7958  
Epoch 38/50  
79/79 [=====] - ETA: 0s - loss: 0.3898 - accuracy: 0.8733  
Epoch 38: val\_loss did not improve from 0.65613  
79/79 [=====] - 42s 528ms/step - loss: 0.3898 - accuracy: 0.8733 - val\_loss: 0.7051 - val\_accuracy: 0.8021  
Epoch 39/50  
79/79 [=====] - ETA: 0s - loss: 0.3820 - accuracy: 0.8760  
Epoch 39: val\_loss did not improve from 0.65613  
79/79 [=====] - 42s 530ms/step - loss: 0.3820 - accuracy: 0.8760 - val\_loss: 0.7427 - val\_accuracy: 0.7900  
Epoch 40/50  
79/79 [=====] - ETA: 0s - loss: 0.3656 - accuracy: 0.8798  
Epoch 40: val\_loss did not improve from 0.65613

79/79 [=====] - 44s 552ms/step - loss: 0.3656 - accuracy: 0.8798 - val\_loss: 0.9921 - val\_accuracy: 0.7436  
Epoch 41/50  
79/79 [=====] - ETA: 0s - loss: 0.3737 - accuracy: 0.8773  
Epoch 41: val\_loss did not improve from 0.65613  
79/79 [=====] - 44s 553ms/step - loss: 0.3737 - accuracy: 0.8773 - val\_loss: 0.8021 - val\_accuracy: 0.7753  
Epoch 42/50  
79/79 [=====] - ETA: 0s - loss: 0.3637 - accuracy: 0.8831  
Epoch 42: val\_loss did not improve from 0.65613  
79/79 [=====] - 43s 537ms/step - loss: 0.3637 - accuracy: 0.8831 - val\_loss: 0.7260 - val\_accuracy: 0.7912  
Epoch 43/50  
79/79 [=====] - ETA: 0s - loss: 0.3668 - accuracy: 0.8810  
Epoch 43: val\_loss did not improve from 0.65613  
79/79 [=====] - 45s 569ms/step - loss: 0.3668 - accuracy: 0.8810 - val\_loss: 0.8774 - val\_accuracy: 0.7559  
Epoch 44/50  
79/79 [=====] - ETA: 0s - loss: 0.3514 - accuracy: 0.8856  
Epoch 44: val\_loss did not improve from 0.65613  
79/79 [=====] - 43s 540ms/step - loss: 0.3514 - accuracy: 0.8856 - val\_loss: 0.8530 - val\_accuracy: 0.7751  
Epoch 45/50  
79/79 [=====] - ETA: 0s - loss: 0.3240 - accuracy: 0.8954  
Epoch 45: val\_loss did not improve from 0.65613  
79/79 [=====] - 45s 571ms/step - loss: 0.3240 - accuracy: 0.8954 - val\_loss: 0.8540 - val\_accuracy: 0.7555  
Epoch 46/50  
79/79 [=====] - ETA: 0s - loss: 0.3412 - accuracy: 0.8902  
Epoch 46: val\_loss did not improve from 0.65613  
79/79 [=====] - 49s 621ms/step - loss: 0.3412 - accuracy: 0.8902 - val\_loss: 0.8199 - val\_accuracy: 0.7790  
Epoch 47/50  
79/79 [=====] - ETA: 0s - loss: 0.3235 - accuracy: 0.8954  
Epoch 47: val\_loss did not improve from 0.65613  
79/79 [=====] - 49s 623ms/step - loss: 0.3235 - accuracy: 0.8954 - val\_loss: 0.8914 - val\_accuracy: 0.7639  
Epoch 48/50  
79/79 [=====] - ETA: 0s - loss: 0.3172 - accuracy: 0.8957  
Epoch 48: val\_loss did not improve from 0.65613  
79/79 [=====] - 52s 657ms/step - loss: 0.3172 - accuracy: 0.8957 - val\_loss: 0.7707 - val\_accuracy: 0.7954  
Epoch 49/50  
79/79 [=====] - ETA: 0s - loss: 0.3117 - accuracy: 0.8976  
Epoch 49: val\_loss did not improve from 0.65613  
79/79 [=====] - 44s 562ms/step - loss: 0.3117 - accuracy: 0.8976 - val\_loss: 0.8442 - val\_accuracy: 0.7903  
Epoch 50/50  
79/79 [=====] - ETA: 0s - loss: 0.3104 - accuracy: 0.8993  
Epoch 50: val\_loss improved from 0.65613 to 0.65273, saving model to best\_model\_augmented.h5  
79/79 [=====] - 51s 645ms/step - loss: 0.3104 - accuracy: 0.8993 - val\_loss: 0.6527 - val\_accuracy: 0.8107  
1250/1250 [=====] - 42s 33ms/step - loss: 0.3486 - accuracy: 0.8863  
313/313 [=====] - 10s 31ms/step - loss: 0.6527 - accuracy: 0.8107

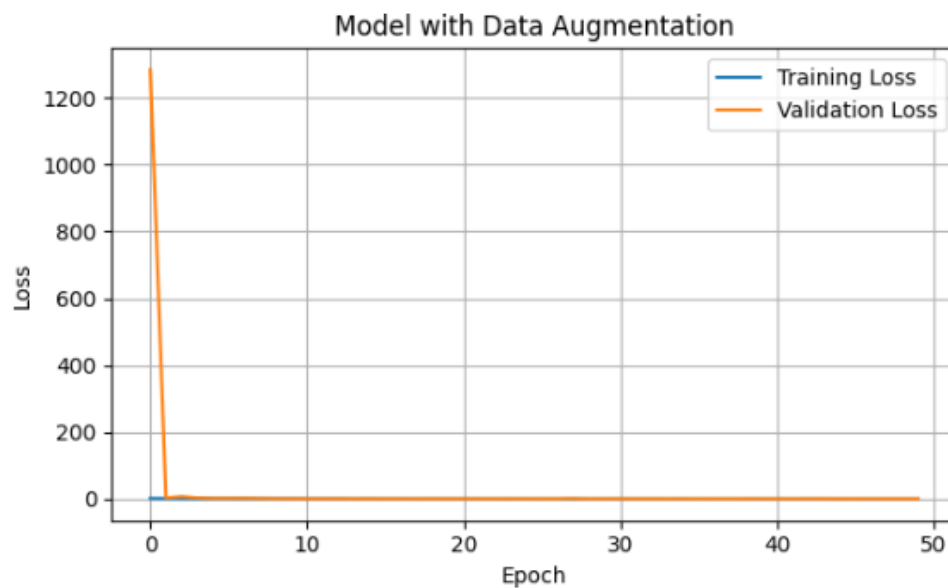
## Plot: Loss vs Epoch (with data augmentation)

```
plt.figure(figsize=(12, 4))

# Plotting Training and Validation Loss
plt.subplot(1, 2, 1)
plt.plot(history_augmented.history['loss'], label='Training Loss', color='blue', marker='o')
plt.plot(history_augmented.history['val_loss'], label='Validation Loss', color='orange', marker='o')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Model with Data Augmentation')
plt.legend()
plt.grid(True)
plt.tight_layout()

plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.title('Training and Validation Loss Over Epochs')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()

plt.show()
```






Plot: Accuracy vs Epoch (with data augmentation)

```
plt.figure(figsize=(12, 4))

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 2)
plt.plot(history_augmented.history['accuracy'], label='Training Accuracy', color='green', marker='o')
plt.plot(history_augmented.history['val_accuracy'], label='Validation Accuracy', color='red', marker='o')
plt.xlabel('Epoch')
plt.ylabel('Accuracy (%)')
plt.title('Model with Data Augmentation')
plt.legend()
plt.grid(True)

plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.title('Training and Validation Accuracy Over Epochs')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)
plt.xlabel('Epoch')
plt.ylabel('Accuracy (%)')
plt.legend()
plt.tight_layout()

plt.show()
```



## ▼ Print training and validation accuracy and loss for the model (with data augmentation)

```
✓ [13] print("\nModel with Data Augmentation:")
0s      print("Training Loss: {:.4f}".format(history_augmented.history['loss'][-1]))
      print("Training Accuracy: {:.4f}%".format(history_augmented.history['accuracy'][-1] * 100))
      print("Validation Loss: {:.4f}".format(history_augmented.history['val_loss'][-1]))
      print("Validation Accuracy: {:.4f}%".format(history_augmented.history['val_accuracy'][-1] * 100))
```

```
Model with Data Augmentation:
Training Loss: 0.3104
Training Accuracy: 89.9325%
Validation Loss: 0.6527
Validation Accuracy: 81.0700%
```

---

## ▼ Print summary of the model (with Data Augmentation)

```
✓ 0s ▶ print("\nModel Summary (with Data Augmentation):")
      model_augmented.summary()
```

---

Model Summary (with Data Augmentation):

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
resnet152v2 (Functional)	(None, 1, 1, 2048)	58331648
conv2d (Conv2D)	(None, 1, 1, 32)	589856
activation (Activation)	(None, 1, 1, 32)	0
conv2d_1 (Conv2D)	(None, 1, 1, 32)	9248
activation_1 (Activation)	(None, 1, 1, 32)	0
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 32)	0
conv2d_2 (Conv2D)	(None, 1, 1, 64)	18496
activation_2 (Activation)	(None, 1, 1, 64)	0
conv2d_3 (Conv2D)	(None, 1, 1, 64)	36928
activation_3 (Activation)	(None, 1, 1, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 64)	0
flatten (Flatten)	(None, 64)	0

dense (Dense)	(None, 512)	33280
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130

```

=====
Total params: 59024586 (225.16 MB)
Trainable params: 58880842 (224.61 MB)
Non-trainable params: 143744 (561.50 KB)

```

## ▼ Evaluate the model with test data

```

✓ [15] test_loss, test_accuracy = model_augmented.evaluate(x_test, y_test)
10s print(f'Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy*100:.2f}%')

313/313 [=====] - 9s 29ms/step - loss: 0.7024 - accuracy: 0.8027
Test Loss: 0.7024, Test Accuracy: 80.27%

```

## ▼ Plotting the confusion matrix as heatmap

```

/ from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
1s

# Predict classes for test set
y_pred = model_augmented.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

# Compute confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred_classes)
print("Confusion Matrix:")
print(conf_matrix)

# Compute classification report
class_report = classification_report(y_true, y_pred_classes, target_names=[str(i) for i in range(10)])
print("Classification Report:")
print(class_report)

# Compute accuracy
accuracy = accuracy_score(y_true, y_pred_classes)
print("Accuracy:", accuracy)

```

```
313/313 [=====] - 8s 26ms/step
```

```

Confusion Matrix:
[[846  14  38   7   9   0   2   3  46  35]
 [ 14 905   0   2   1   0   9   0  17  52]
 [ 36   6 774  21  48   8  57  28  11  11]
 [ 39  17  63 560  66  58  69  65  24  39]
 [ 14   4  47  14 800   8  40  53  10  10]
 [ 19   9  42 162  57 542  39 100  10  20]
 [  9  12  25  21  10   1 895   8   6  13]
 [ 21   5  12   5  22   8   3 907   2  15]
 [ 41  19   3   1   2   0   4   0 908  22]
 [ 22  61   1   3   1   0   1   5  16 890]]

```

### Classification Report:

	precision	recall	f1-score	support
0	0.80	0.85	0.82	1000
1	0.86	0.91	0.88	1000
2	0.77	0.77	0.77	1000
3	0.70	0.56	0.62	1000
4	0.79	0.80	0.79	1000
5	0.87	0.54	0.67	1000
6	0.80	0.90	0.84	1000
7	0.78	0.91	0.84	1000
8	0.86	0.91	0.89	1000
9	0.80	0.89	0.84	1000
accuracy			0.80	10000
macro avg	0.80	0.80	0.80	10000
weighted avg	0.80	0.80	0.80	10000


Accuracy: 0.8027

### Confusion matrix plot

```
[20] import seaborn as sns
import matplotlib.pyplot as plt

# Compute confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred_classes)

# Plot confusion matrix heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=[str(i) for i in range(10)], yticklabels=[str(i) for i in range(10)])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



---

Note: Due to technical issue in google colab output was not visible. So, needed to rerun the code for trail 1.

## Trail 2: Code

**Google Colab link for trail 1 code:**

[https://colab.research.google.com/drive/1UzXngwYTYcR31BoRGX478JfiOddh\\_cGf?usp=sharing](https://colab.research.google.com/drive/1UzXngwYTYcR31BoRGX478JfiOddh_cGf?usp=sharing)

```
[ ] # Import necessary libraries
import numpy as np
from keras.datasets import cifar10
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, Activation
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt

# Load CIFAR-10 dataset and split into training and testing sets
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Create a validation set by randomly selecting 20% of the training images
validation_split = 0.2
validation_samples = int(len(x_train) * validation_split)
x_val = x_train[:validation_samples]
y_val = y_train[:validation_samples]
x_train = x_train[validation_samples:]
y_train = y_train[validation_samples:]

# Scale pixel values to a range between 0 and 1
x_train = x_train.astype('float32') / 255
x_val = x_val.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# Convert labels to binary class matrices
y_train = to_categorical(y_train, 10)
y_val = to_categorical(y_val, 10)
y_test = to_categorical(y_test, 10)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 6s 0us/step

```

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers import Dense, Flatten, Activation
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from keras.preprocessing.image import ImageDataGenerator
from keras.src.layers.serialization import activation

# Build CNN model with data augmentation
data_generator = ImageDataGenerator(rotation_range=10, width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
data_generator.fit(x_train)

model_augmented = Sequential()
pretraining_model = tf.keras.applications.ResNet152V2(
    include_top=False,
    weights="imagenet",
    input_tensor=None,
    input_shape=(32, 32, 3),
    pooling=None, # No global pooling in the pre-trained model
    classes=10
)

for layer in pretraining_model.layers:
    layer.trainable = True

```

```

model_augmented.add(pretraining_model)
model_augmented.add(Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(32, 32, 3)))
model_augmented.add(Activation('relu'))
model_augmented.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model_augmented.add(Activation('relu'))
model_augmented.add(MaxPooling2D(pool_size=(1, 1)))
model_augmented.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model_augmented.add(Activation('relu'))
model_augmented.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model_augmented.add(Activation('relu'))
model_augmented.add(MaxPooling2D(pool_size=(1, 1)))
model_augmented.add(Flatten())
model_augmented.add(Dense(512, activation='relu'))
model_augmented.add(Dropout(0.5))
model_augmented.add(Dense(10, activation='softmax'))

model_augmented.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.001), metrics=['accuracy'])
checkpoint_augmented = ModelCheckpoint('best_model_augmented.h5', save_best_only=True, monitor='val_loss', mode='min', verbose=1)
history_augmented = model_augmented.fit(data_generator.flow(x_train, y_train, batch_size=32), epochs=50, validation_data=(x_val, y_val), callbacks=[checkpoint_augmented])
model_augmented.load_weights('best_model_augmented.h5')
train_loss_augmented, train_accuracy_augmented = model_augmented.evaluate(x_train, y_train)
val_loss_augmented, val_accuracy_augmented = model_augmented.evaluate(x_val, y_val)

```

```

# Step f: Plot training and validation loss and accuracy for the second model (with data augmentation)
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history_augmented.history['loss'], label='Training Loss')
plt.plot(history_augmented.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Model with Data Augmentation')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(history_augmented.history['accuracy'], label='Training Accuracy')
plt.plot(history_augmented.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy (%)')
plt.title('Model with Data Augmentation')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

```

```

# Step f: Print training and validation accuracy and loss for the second model (with data augmentation)
print("\nModel with Data Augmentation:")
print("Training Loss: {:.4f}".format(history_augmented.history['loss'][-1]))
print("Training Accuracy: {:.4f}%".format(history_augmented.history['accuracy'][-1] * 100))
print("Validation Loss: {:.4f}".format(history_augmented.history['val_loss'][-1]))
print("Validation Accuracy: {:.4f}%".format(history_augmented.history['val_accuracy'][-1] * 100))

# Step f: Print summary of the second model (with Data Augmentation)
print("\nModel Summary (with Data Augmentation):")
model_augmented.summary()

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5)

234545216/234545216 [=====] - 7s 0us/step

WARNING:absl:lr is deprecated in Keras optimizer, please use

`learning\_rate` or use the legacy optimizer,

e.g., tf.keras.optimizers.legacy.Adam.

Epoch 1/50

1250/1250 [=====] - ETA: 0s - loss: 2.1903 - accuracy: 0.1489

Epoch 1: val\_loss improved from inf to 2.17104, saving model to best\_model\_augmented.h5

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3079:

UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native

Keras format, e.g. `model.save('my\_model.keras')`. saving\_api.save\_model(

1250/1250 [=====] - 275s 146ms/step - loss: 2.1903 - accuracy: 0.1489 - val\_loss: 2.1710 - val\_accuracy: 0.1583

Epoch 2/50

1250/1250 [=====] - ETA: 0s - loss: 1.9628 - accuracy: 0.2121

Epoch 2: val\_loss improved from 2.17104 to 1.87865, saving model to best\_model\_augmented.h5

1250/1250 [=====] - 179s 143ms/step - loss: 1.9628 - accuracy: 0.2121 - val\_loss: 1.8786 - val\_accuracy: 0.2542

Epoch 3/50

1250/1250 [=====] - ETA: 0s - loss: 1.8580 - accuracy: 0.2725

Epoch 3: val\_loss improved from 1.87865 to 1.79871, saving model to best\_model\_augmented.h5

1250/1250 [=====] - 185s 148ms/step - loss: 1.8580 - accuracy: 0.2725 - val\_loss: 1.7987 - val\_accuracy: 0.2711

Epoch 4/50

1250/1250 [=====] - ETA: 0s - loss: 1.7364 - accuracy: 0.3179

Epoch 4: val\_loss did not improve from 1.79871

1250/1250 [=====] - 174s 139ms/step - loss: 1.7364 - accuracy: 0.3179 - val\_loss: 2.1845 - val\_accuracy: 0.2643

Epoch 5/50

1250/1250 [=====] - ETA: 0s - loss: 1.6791 - accuracy: 0.3440

Epoch 5: val\_loss did not improve from 1.79871

1250/1250 [=====] - 171s 137ms/step - loss: 1.6791 - accuracy: 0.3440 - val\_loss: 1.9083 - val\_accuracy: 0.3516

Epoch 6/50

1250/1250 [=====] - ETA: 0s - loss: 1.5426 - accuracy: 0.4062



Epoch 6: val\_loss improved from 1.79871 to 1.48656, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 183s 147ms/step - loss: 1.5426 -  
accuracy: 0.4062 - val\_loss: 1.4866 - val\_accuracy: 0.4409  
Epoch 7/50  
1250/1250 [=====] - ETA: 0s - loss: 1.4258 -  
accuracy: 0.4710  
Epoch 7: val\_loss improved from 1.48656 to 1.26602, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 188s 150ms/step - loss: 1.4258 -  
accuracy: 0.4710 - val\_loss: 1.2660 - val\_accuracy: 0.5384  
Epoch 8/50  
1250/1250 [=====] - ETA: 0s - loss: 1.3222 -  
accuracy: 0.5253  
Epoch 8: val\_loss did not improve from 1.26602  
1250/1250 [=====] - 176s 141ms/step - loss: 1.3222 -  
accuracy: 0.5253 - val\_loss: 1.7111 - val\_accuracy: 0.4583  
Epoch 9/50  
1250/1250 [=====] - ETA: 0s - loss: 1.2560 -  
accuracy: 0.5548  
Epoch 9: val\_loss did not improve from 1.26602  
1250/1250 [=====] - 176s 141ms/step - loss: 1.2560 -  
accuracy: 0.5548 - val\_loss: 1.3473 - val\_accuracy: 0.5447  
Epoch 10/50  
1250/1250 [=====] - ETA: 0s - loss: 1.1941 -  
accuracy: 0.5836  
Epoch 10: val\_loss did not improve from 1.26602  
1250/1250 [=====] - 177s 142ms/step - loss: 1.1941 -  
accuracy: 0.5836 - val\_loss: 1.4566 - val\_accuracy: 0.5001  
Epoch 11/50  
1250/1250 [=====] - ETA: 0s - loss: 1.1316 -  
accuracy: 0.6068  
Epoch 11: val\_loss improved from 1.26602 to 1.05279, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 191s 153ms/step - loss: 1.1316 -  
accuracy: 0.6068 - val\_loss: 1.0528 - val\_accuracy: 0.6223  
Epoch 12/50  
1250/1250 [=====] - ETA: 0s - loss: 1.0900 -  
accuracy: 0.6267  
Epoch 12: val\_loss did not improve from 1.05279  
1250/1250 [=====] - 176s 141ms/step - loss: 1.0900 -  
accuracy: 0.6267 - val\_loss: 1.1745 - val\_accuracy: 0.5988  
Epoch 13/50  
1250/1250 [=====] - ETA: 0s - loss: 1.0437 -  
accuracy: 0.6434  
Epoch 13: val\_loss improved from 1.05279 to 0.96482, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 193s 154ms/step - loss: 1.0437 -  
accuracy: 0.6434 - val\_loss: 0.9648 - val\_accuracy: 0.6687  
Epoch 14/50  
1250/1250 [=====] - ETA: 0s - loss: 1.0189 -  
accuracy: 0.6503  
Epoch 14: val\_loss did not improve from 0.96482  
1250/1250 [=====] - 175s 140ms/step - loss: 1.0189 -  
accuracy: 0.6503 - val\_loss: 1.0367 - val\_accuracy: 0.6449  
Epoch 15/50


1250/1250 [=====] - ETA: 0s - loss: 0.9789 -  
accuracy: 0.6647  
Epoch 15: val\_loss improved from 0.96482 to 0.86103, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 191s 153ms/step - loss: 0.9789 -  
accuracy: 0.6647 - val\_loss: 0.8610 - val\_accuracy: 0.7066  
Epoch 16/50  
1250/1250 [=====] - ETA: 0s - loss: 0.9473 -  
accuracy: 0.6762  
Epoch 16: val\_loss improved from 0.86103 to 0.85211, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 189s 151ms/step - loss: 0.9473 -  
accuracy: 0.6762 - val\_loss: 0.8521 - val\_accuracy: 0.7065  
Epoch 17/50  
1250/1250 [=====] - ETA: 0s - loss: 0.9253 -  
accuracy: 0.6855  
Epoch 17: val\_loss did not improve from 0.85211  
1250/1250 [=====] - 174s 139ms/step - loss: 0.9253 -  
accuracy: 0.6855 - val\_loss: 0.8540 - val\_accuracy: 0.7126  
Epoch 18/50  
1250/1250 [=====] - ETA: 0s - loss: 0.8930 -  
accuracy: 0.6974  
Epoch 18: val\_loss improved from 0.85211 to 0.81993, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 189s 151ms/step - loss: 0.8930 -  
accuracy: 0.6974 - val\_loss: 0.8199 - val\_accuracy: 0.7252  
Epoch 19/50  
1250/1250 [=====] - ETA: 0s - loss: 0.8820 -  
accuracy: 0.7049  
Epoch 19: val\_loss did not improve from 0.81993  
1250/1250 [=====] - 176s 141ms/step - loss: 0.8820 -  
accuracy: 0.7049 - val\_loss: 1.0898 - val\_accuracy: 0.6573  
Epoch 20/50  
1250/1250 [=====] - ETA: 0s - loss: 0.8472 -  
accuracy: 0.7145  
Epoch 20: val\_loss did not improve from 0.81993  
1250/1250 [=====] - 174s 139ms/step - loss: 0.8472 -  
accuracy: 0.7145 - val\_loss: 1.0834 - val\_accuracy: 0.6457  
Epoch 21/50  
1250/1250 [=====] - ETA: 0s - loss: 0.8377 -  
accuracy: 0.7216  
Epoch 21: val\_loss did not improve from 0.81993  
1250/1250 [=====] - 177s 141ms/step - loss: 0.8377 -  
accuracy: 0.7216 - val\_loss: 0.8673 - val\_accuracy: 0.7164  
Epoch 22/50  
1250/1250 [=====] - ETA: 0s - loss: 0.8109 -  
accuracy: 0.7333  
Epoch 22: val\_loss improved from 0.81993 to 0.78372, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 189s 151ms/step - loss: 0.8109 -  
accuracy: 0.7333 - val\_loss: 0.7837 - val\_accuracy: 0.7403  
Epoch 23/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7911 -  
accuracy: 0.7395  
Epoch 23: val\_loss did not improve from 0.78372  
1250/1250 [=====] - 173s 139ms/step - loss: 0.7911 -  
accuracy: 0.7395 - val\_loss: 0.8516 - val\_accuracy: 0.7195

Epoch 24/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7868 -  
accuracy: 0.7405  
Epoch 24: val\_loss did not improve from 0.78372  
1250/1250 [=====] - 174s 139ms/step - loss: 0.7868 -  
accuracy: 0.7405 - val\_loss: 0.9313 - val\_accuracy: 0.6943  
Epoch 25/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7619 -  
accuracy: 0.7498  
Epoch 25: val\_loss did not improve from 0.78372  
1250/1250 [=====] - 173s 138ms/step - loss: 0.7619 -  
accuracy: 0.7498 - val\_loss: 0.9183 - val\_accuracy: 0.7058  
Epoch 26/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7463 -  
accuracy: 0.7565  
Epoch 26: val\_loss did not improve from 0.78372  
1250/1250 [=====] - 174s 139ms/step - loss: 0.7463 -  
accuracy: 0.7565 - val\_loss: 0.8697 - val\_accuracy: 0.7116  
Epoch 27/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7253 -  
accuracy: 0.7608  
Epoch 27: val\_loss improved from 0.78372 to 0.73048, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 189s 151ms/step - loss: 0.7253 -  
accuracy: 0.7608 - val\_loss: 0.7305 - val\_accuracy: 0.7601  
Epoch 28/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7182 -  
accuracy: 0.7648  
Epoch 28: val\_loss did not improve from 0.73048  
1250/1250 [=====] - 174s 139ms/step - loss: 0.7182 -  
accuracy: 0.7648 - val\_loss: 1.0052 - val\_accuracy: 0.7043  
Epoch 29/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7059 -  
accuracy: 0.7683  
Epoch 29: val\_loss improved from 0.73048 to 0.70456, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 189s 151ms/step - loss: 0.7059 -  
accuracy: 0.7683 - val\_loss: 0.7046 - val\_accuracy: 0.7739  
Epoch 30/50  
1250/1250 [=====] - ETA: 0s - loss: 0.7003 -  
accuracy: 0.7728  
Epoch 30: val\_loss did not improve from 0.70456  
1250/1250 [=====] - 173s 139ms/step - loss: 0.7003 -  
accuracy: 0.7728 - val\_loss: 0.7929 - val\_accuracy: 0.7368  
Epoch 31/50  
1250/1250 [=====] - ETA: 0s - loss: 0.6795 -  
accuracy: 0.7785  
Epoch 31: val\_loss did not improve from 0.70456  
1250/1250 [=====] - 174s 139ms/step - loss: 0.6795 -  
accuracy: 0.7785 - val\_loss: 1.2060 - val\_accuracy: 0.6457  
Epoch 32/50  
1250/1250 [=====] - ETA: 0s - loss: 0.6672 -  
accuracy: 0.7848  
Epoch 32: val\_loss did not improve from 0.70456  
1250/1250 [=====] - 171s 137ms/step - loss: 0.6672 -  
accuracy: 0.7848 - val\_loss: 0.8987 - val\_accuracy: 0.7109  
Epoch 33/50

1250/1250 [=====] - ETA: 0s - loss: 0.6516 -  
accuracy: 0.7869  
Epoch 33: val\_loss did not improve from 0.70456  
1250/1250 [=====] - 173s 139ms/step - loss: 0.6516 -  
accuracy: 0.7869 - val\_loss: 0.7901 - val\_accuracy: 0.7387  
Epoch 34/50  
1250/1250 [=====] - ETA: 0s - loss: 0.6473 -  
accuracy: 0.7892  
Epoch 34: val\_loss did not improve from 0.70456  
1250/1250 [=====] - 171s 137ms/step - loss: 0.6473 -  
accuracy: 0.7892 - val\_loss: 0.8832 - val\_accuracy: 0.7181  
Epoch 35/50  
1250/1250 [=====] - ETA: 0s - loss: 0.6313 -  
accuracy: 0.7965  
Epoch 35: val\_loss did not improve from 0.70456  
1250/1250 [=====] - 174s 139ms/step - loss: 0.6313 -  
accuracy: 0.7965 - val\_loss: 0.7414 - val\_accuracy: 0.7606  
Epoch 36/50  
1250/1250 [=====] - ETA: 0s - loss: 0.6233 -  
accuracy: 0.7972  
Epoch 36: val\_loss improved from 0.70456 to 0.69908, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 192s 153ms/step - loss: 0.6233 -  
accuracy: 0.7972 - val\_loss: 0.6991 - val\_accuracy: 0.7723  
Epoch 37/50  
1250/1250 [=====] - ETA: 0s - loss: 0.6128 -  
accuracy: 0.8026  
Epoch 37: val\_loss did not improve from 0.69908  
1250/1250 [=====] - 173s 139ms/step - loss: 0.6128 -  
accuracy: 0.8026 - val\_loss: 1.0388 - val\_accuracy: 0.6900  
Epoch 38/50  
1250/1250 [=====] - ETA: 0s - loss: 0.5998 -  
accuracy: 0.8072  
Epoch 38: val\_loss did not improve from 0.69908  
1250/1250 [=====] - 174s 139ms/step - loss: 0.5998 -  
accuracy: 0.8072 - val\_loss: 0.7043 - val\_accuracy: 0.7727  
Epoch 39/50  
1250/1250 [=====] - ETA: 0s - loss: 0.5868 -  
accuracy: 0.8096  
Epoch 39: val\_loss did not improve from 0.69908  
1250/1250 [=====] - 177s 142ms/step - loss: 0.5868 -  
accuracy: 0.8096 - val\_loss: 0.8743 - val\_accuracy: 0.7303  
Epoch 40/50  
1250/1250 [=====] - ETA: 0s - loss: 0.5904 -  
accuracy: 0.8077  
Epoch 40: val\_loss did not improve from 0.69908  
1250/1250 [=====] - 177s 142ms/step - loss: 0.5904 -  
accuracy: 0.8077 - val\_loss: 1.2257 - val\_accuracy: 0.6585  
Epoch 41/50  
1250/1250 [=====] - ETA: 0s - loss: 0.5761 -  
accuracy: 0.8126  
Epoch 41: val\_loss improved from 0.69908 to 0.63643, saving model to  
best\_model\_augmented.h5  
1250/1250 [=====] - 188s 150ms/step - loss: 0.5761 -  
accuracy: 0.8126 - val\_loss: 0.6364 - val\_accuracy: 0.7912  
Epoch 42/50

```
1250/1250 [=====] - ETA: 0s - loss: 0.5689 -
accuracy: 0.8169
Epoch 42: val_loss did not improve from 0.63643
1250/1250 [=====] - 178s 142ms/step - loss: 0.5689 -
accuracy: 0.8169 - val_loss: 0.7956 - val_accuracy: 0.7417
Epoch 43/50
1250/1250 [=====] - ETA: 0s - loss: 0.5629 -
accuracy: 0.8215
Epoch 43: val_loss did not improve from 0.63643
1250/1250 [=====] - 174s 139ms/step - loss: 0.5629 -
accuracy: 0.8215 - val_loss: 0.7473 - val_accuracy: 0.7579
Epoch 44/50
1250/1250 [=====] - ETA: 0s - loss: 0.5514 -
accuracy: 0.8208
Epoch 44: val_loss did not improve from 0.63643
1250/1250 [=====] - 172s 138ms/step - loss: 0.5514 -
accuracy: 0.8208 - val_loss: 0.8438 - val_accuracy: 0.7454
Epoch 45/50
1250/1250 [=====] - ETA: 0s - loss: 0.5476 -
accuracy: 0.8219
Epoch 45: val_loss did not improve from 0.63643
1250/1250 [=====] - 171s 136ms/step - loss: 0.5476 -
accuracy: 0.8219 - val_loss: 0.6774 - val_accuracy: 0.7862
Epoch 46/50
1250/1250 [=====] - ETA: 0s - loss: 0.5293 -
accuracy: 0.8300
Epoch 46: val_loss did not improve from 0.63643
1250/1250 [=====] - 173s 138ms/step - loss: 0.5293 -
accuracy: 0.8300 - val_loss: 1.2502 - val_accuracy: 0.6492
Epoch 47/50
1250/1250 [=====] - ETA: 0s - loss: 0.5354 -
accuracy: 0.8306
Epoch 47: val_loss did not improve from 0.63643
1250/1250 [=====] - 171s 137ms/step - loss: 0.5354 -
accuracy: 0.8306 - val_loss: 1.5224 - val_accuracy: 0.5833
Epoch 48/50
1250/1250 [=====] - ETA: 0s - loss: 0.5241 -
accuracy: 0.8313
Epoch 48: val_loss improved from 0.63643 to 0.63485, saving model to
best_model_augmented.h5
1250/1250 [=====] - 184s 147ms/step - loss: 0.5241 -
accuracy: 0.8313 - val_loss: 0.6349 - val_accuracy: 0.7962
Epoch 49/50
1250/1250 [=====] - ETA: 0s - loss: 0.5212 -
accuracy: 0.8338
Epoch 49: val_loss did not improve from 0.63485
1250/1250 [=====] - 170s 136ms/step - loss: 0.5212 -
accuracy: 0.8338 - val_loss: 0.7516 - val_accuracy: 0.7586
Epoch 50/50
1250/1250 [=====] - ETA: 0s - loss: 0.5060 -
accuracy: 0.8378
Epoch 50: val_loss did not improve from 0.63485
1250/1250 [=====] - 174s 139ms/step - loss: 0.5060 -
accuracy: 0.8378 - val_loss: 0.7224 - val_accuracy: 0.7666
1250/1250 [=====] - 39s 31ms/step - loss: 0.4508 -
accuracy: 0.8487
```

313/313 [=====] - 10s 33ms/step - loss: 0.6349 - accuracy: 0.7962



Model with Data Augmentation:  
Training Loss: 0.5060  
Training Accuracy: 83.7750%  
Validation Loss: 0.7224  
Validation Accuracy: 76.6600%

Model Summary (with Data Augmentation):  
Model: "sequential\_1"

Layer (type)	Output Shape	Param #
resnet152v2 (Functional)	(None, 1, 1, 2048)	58331648
conv2d (Conv2D)	(None, 1, 1, 32)	589856
activation (Activation)	(None, 1, 1, 32)	0
conv2d_1 (Conv2D)	(None, 1, 1, 32)	9248
activation_1 (Activation)	(None, 1, 1, 32)	0
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 32)	0
conv2d_2 (Conv2D)	(None, 1, 1, 64)	18496
activation_2 (Activation)	(None, 1, 1, 64)	0
conv2d_3 (Conv2D)	(None, 1, 1, 64)	36928
activation_3 (Activation)	(None, 1, 1, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 64)	0
flatten (Flatten)	(None, 64)	0
dense (Dense)	(None, 512)	33280
dropout (Dropout)	(None, 512)	0

dense_1 (Dense)	(None, 10)	5130
-----------------	------------	------

```
=====
Total params: 59024586 (225.16 MB)
Trainable params: 58880842 (224.61 MB)
Non-trainable params: 143744 (561.50 KB)
=====
```

```
[ ] # Evaluate the model with test data
    test_loss, test_accuracy = model_augmented.evaluate(x_test, y_test)
    print(f'Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy*100:.2f}%')

313/313 [=====] - 10s 32ms/step - loss: 0.6532 - accuracy: 0.7895
Test Loss: 0.6532, Test Accuracy: 78.95%
```